

A1

```
1)
#include <iostream>
#include <unistd.h>    // for pipe, fork, read, write
#include <sys/types.h>
#include <sys/wait.h>  // for wait

int main() {
    int pipe1[2], pipe2[2];
    // Create two pipes: pipe1 for parent->child, pipe2 for child->parent
    if (pipe(pipe1) == -1 || pipe(pipe2) == -1) {
        std::cerr << "Pipe creation failed\n";
        return 1;
    }

    pid_t pid = fork();
    if (pid < 0) {
        std::cerr << "Fork failed\n";
        return 1;
    } else if (pid > 0) {
        // Parent process
        close(pipe1[0]); // Close unused read end
        close(pipe2[1]); // Close unused write end

        char input[100];
        std::cout << "Enter a string: ";
        std::cin.getline(input, 100);

        // Compute length manually (no string functions)
        int len = 0;
        while (input[len] != '\0') {
            len++;
        }

        // Send string (including null terminator) to child
        write(pipe1[1], input, len + 1);

        // Read concatenated result from child
        char result[200];
        read(pipe2[0], result, sizeof(result));
        std::cout << "Concatenated string: " << result << std::endl;

        close(pipe1[1]);
        close(pipe2[0]);
        wait(NULL); // Wait for child to finish
    } else {
        // Child process
        close(pipe1[1]); // Close unused write end
        close(pipe2[0]); // Close unused read end

        // Read input from parent
        char received[100];
        read(pipe1[0], received, sizeof(received));

        // String to concatenate
        char extra[] = " - concatenated by child";

        // Manual concatenation
        char result[200];
```

```

    int i = 0;
    while (received[i] != '\0') {
        result[i] = received[i];
        i++;
    }
    int j = 0;
    while (extra[j] != '\0') {
        result[i + j] = extra[j];
        j++;
    }
    result[i + j] = '\0';

    // Send back to parent
    write(pipe2[1], result, i + j + 1);

    close(pipe1[0]);
    close(pipe2[1]);
}

return 0;
}

2)
#include <iostream>
#include <unistd.h>    // for pipe, fork, read, write
#include <sys/types.h>
#include <sys/wait.h>  // for wait
#include <cstdlib>      // for exit

int main() {
    int pipe1[2], pipe2[2];
    // pipe1: parent -> child, pipe2: child -> parent
    if (pipe(pipe1) == -1 || pipe(pipe2) == -1) {
        std::cerr << "Failed to create pipes\n";
        return 1;
    }

    pid_t pid = fork();
    if (pid < 0) {
        std::cerr << "Fork failed\n";
        return 1;
    }

    if (pid > 0) {
        // Parent process
        close(pipe1[0]); // close read end of pipe1
        close(pipe2[1]); // close write end of pipe2

        int rows, cols;
        std::cout << "Enter number of rows: ";
        std::cin >> rows;
        std::cout << "Enter number of columns: ";
        std::cin >> cols;

        int size = rows * cols;
        int* A = new int[size];
        int* B = new int[size];

        std::cout << "Enter elements of first matrix:\n";
        for (int i = 0; i < size; ++i) {

```

```

        std::cin >> A[i];
    }
    std::cout << "Enter elements of second matrix:\n";
    for (int i = 0; i < size; ++i) {
        std::cin >> B[i];
    }

    // Send rows and cols
    write(pipe1[1], &rows, sizeof(rows));
    write(pipe1[1], &cols, sizeof(cols));
    // Send matrix data
    write(pipe1[1], A, size * sizeof(int));
    write(pipe1[1], B, size * sizeof(int));

    delete[] A;
    delete[] B;
    close(pipe1[1]);

    // Read result dims (optional)
    int r2, c2;
    read(pipe2[0], &r2, sizeof(r2));
    read(pipe2[0], &c2, sizeof(c2));
    int size2 = r2 * c2;
    int* C = new int[size2];
    read(pipe2[0], C, size2 * sizeof(int));

    std::cout << "Sum matrix is:\n";
    for (int i = 0; i < r2; ++i) {
        for (int j = 0; j < c2; ++j) {
            std::cout << C[i * c2 + j] << " ";
        }
        std::cout << "\n";
    }

    delete[] C;
    close(pipe2[0]);
    wait(NULL);
} else {
    // Child process
    close(pipe1[1]); // close write end of pipe1
    close(pipe2[0]); // close read end of pipe2

    int rows, cols;
    // Read dims
    read(pipe1[0], &rows, sizeof(rows));
    read(pipe1[0], &cols, sizeof(cols));
    int size = rows * cols;
    int* A = new int[size];
    int* B = new int[size];
    // Read matrices
    read(pipe1[0], A, size * sizeof(int));
    read(pipe1[0], B, size * sizeof(int));
    close(pipe1[0]);

    int* C = new int[size];
    for (int i = 0; i < size; ++i) {
        C[i] = A[i] + B[i];
    }

    // Send back dims and summed matrix

```

```

write(pipe2[1], &rows, sizeof(rows));
write(pipe2[1], &cols, sizeof(cols));
write(pipe2[1], C, size * sizeof(int));

delete[] A;
delete[] B;
delete[] C;
close(pipe2[1]);
}

return 0;
}

```

A2

2) server

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <semaphore.h>

// Shared data layout
typedef struct {
    sem_t sem; // POSIX unnamed semaphore
    int counter; // shared counter
} shared_data_t;

int main() {
    const char *path = "/tmp/my_shm";
    int fd = open(path, O_RDWR | O_CREAT, 0666);
    if (fd == -1) { perror("open"); exit(EXIT_FAILURE); }
    if (ftruncate(fd, sizeof(shared_data_t)) == -1) { perror("ftruncate"); exit(EXIT_FAILURE); }

    shared_data_t *shm = mmap(
        NULL,
        sizeof(shared_data_t),
        PROT_READ | PROT_WRITE,
        MAP_SHARED,
        fd,
        0
    );
    if (shm == MAP_FAILED) { perror("mmap"); exit(EXIT_FAILURE); }

    // Initialize semaphore for inter-process use, and counter
    sem_init(&shm->sem, 1, 1);
    shm->counter = 0;

    printf("[Server] Ready. Press Enter to increment counter, 'q'+Enter to quit.\n");
    char buf[8];
    while (fgets(buf, sizeof(buf), stdin)) {
        if (buf[0] == 'q') break;

        // Critical section
        sem_wait(&shm->sem);
        shm->counter++;
        printf("[Server] Counter = %d\n", shm->counter);
    }
}

```

```

    sem_post(&shm->sem);
}

// Cleanup
sem_destroy(&shm->sem);
munmap(shm, sizeof(shared_data_t));
close(fd);
printf("[Server] Exiting.\n");
return 0;
}

```

3) client

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <semaphore.h>

```

```

typedef struct {
    sem_t sem;
    int counter;
} shared_data_t;

```

```

int main() {
    const char *path = "/tmp/my_shm";
    int fd = open(path, O_RDWR);
    if (fd == -1) { perror("open"); exit(EXIT_FAILURE); }

    shared_data_t *shm = mmap(
        NULL,
        sizeof(shared_data_t),
        PROT_READ | PROT_WRITE,
        MAP_SHARED,
        fd,
        0
    );
    if (shm == MAP_FAILED) { perror("mmap"); exit(EXIT_FAILURE); }

    printf("[Client] Ready. Press Enter to read counter, 'q'+Enter to quit.\n");
    char buf[8];
    while (fgets(buf, sizeof(buf), stdin)) {
        if (buf[0] == 'q') break;

        sem_wait(&shm->sem);
        printf("[Client] Counter = %d\n", shm->counter);
        sem_post(&shm->sem);
    }

    munmap(shm, sizeof(shared_data_t));
    close(fd);
    printf("[Client] Exiting.\n");
    return 0;
}

```

1)doubt

A3

```

1)server
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 1024

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <port>\n", argv[0]);
        exit(EXIT_FAILURE);
    }
    int port = atoi(argv[1]);

    // 1) Create socket
    int sock = socket(AF_INET, SOCK_DGRAM, 0);
    if (sock < 0) {
        perror("socket");
        exit(EXIT_FAILURE);
    }

    // 2) Bind to all interfaces on given port
    struct sockaddr_in srv_addr = {0};
    srv_addr.sin_family = AF_INET;
    srv_addr.sin_addr.s_addr = INADDR_ANY;
    srv_addr.sin_port = htons(port);

    if (bind(sock, (struct sockaddr*)&srv_addr, sizeof(srv_addr)) < 0) {
        perror("bind");
        close(sock);
        exit(EXIT_FAILURE);
    }

    printf("UDP server listening on port %d\n", port);

    // 3) Receive loop
    char buf[BUF_SIZE];
    struct sockaddr_in cli_addr;
    socklen_t cli_len = sizeof(cli_addr);

    while (1) {
        ssize_t n = recvfrom(sock, buf, BUF_SIZE-1, 0,
                             (struct sockaddr*)&cli_addr, &cli_len);
        if (n < 0) {
            perror("recvfrom");
            break;
        }
        buf[n] = '\0';
        printf("Received from %s:%d: \"%s\"\n",
              inet_ntoa(cli_addr.sin_addr),
              ntohs(cli_addr.sin_port),
              buf);
        // (Optionally) send a reply:
        // sendto(sock, buf, n, 0, (struct sockaddr*)&cli_addr, cli_len);
    }

    close(sock);
}

```

```
    return 0;
}
```

client

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>
```

```
#define BUF_SIZE 1024
```

```
int main(int argc, char *argv[]) {
    if (argc != 4) {
        fprintf(stderr, "Usage: %s <server-ip> <server-port> <message>\n", argv[0]);
        exit(EXIT_FAILURE);
    }
    const char *srv_ip = argv[1];
    int      srv_port = atoi(argv[2]);
    const char *msg = argv[3];

    // 1) Create socket
    int sock = socket(AF_INET, SOCK_DGRAM, 0);
    if (sock < 0) {
        perror("socket");
        exit(EXIT_FAILURE);
    }

    // 2) Fill in server address
    struct sockaddr_in srv_addr = {0};
    srv_addr.sin_family = AF_INET;
    srv_addr.sin_port = htons(srv_port);
    if (inet_aton(srv_ip, &srv_addr.sin_addr) == 0) {
        fprintf(stderr, "Invalid IP address: %s\n", srv_ip);
        close(sock);
        exit(EXIT_FAILURE);
    }

    // 3) Send the message
    ssize_t sent = sendto(sock, msg, strlen(msg), 0,
                          (struct sockaddr*)&srv_addr, sizeof(srv_addr));
    if (sent < 0) {
        perror("sendto");
        close(sock);
        exit(EXIT_FAILURE);
    }

    printf("Sent \"%s\" to %s:%d\n", msg, srv_ip, srv_port);

    // (Optionally) receive a reply:
    // char buf[BUF_SIZE];
    // struct sockaddr_in from;
    // socklen_t fromlen = sizeof(from);
    // ssize_t n = recvfrom(sock, buf, BUF_SIZE-1, 0, (struct sockaddr*)&from, &fromlen);
    // if (n>0) { buf[n]='\0'; printf("Reply: %s\n", buf); }

    close(sock);
    return 0;
}
```

```
}
```

```
2)
```

```
// file: mpi_master_worker.c
```

```
#include <mpi.h>
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
#define TAG_STOP 0
```

```
#define TAG_TASK 1
```

```
#define TAG_RESULT 2
```

```
// Define your task and result payloads:
```

```
typedef struct {  
    long start;    // inclusive  
    long end;      // exclusive  
} task_t;
```

```
typedef struct {  
    double partial_sum;  
    long start;  
    long end;  
} result_t;
```

```
// Example heavy-compute function
```

```
double f(long i) {  
    // Just a dummy; replace with your real work.  
    return 1.0 / (1.0 + i*i);  
}
```

```
int main(int argc, char** argv) {  
    MPI_Init(&argc, &argv);
```

```
    int rank, P;  
    MPI_Comm_size(MPI_COMM_WORLD, &P);  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    const long N = 100000000L;    // total elements  
    const long CHUNK_SIZE = 100000; // elements per task
```

```
    if (rank == 0) {  
        // ----- Master -----  
        long num_tasks = (N + CHUNK_SIZE - 1) / CHUNK_SIZE;  
        long next_task = 0;  
        double total = 0.0;  
  
        // 1) Prime the workers with one task each:  
        for (int worker = 1; worker < P && next_task < num_tasks; worker++) {  
            task_t t;  
            t.start = next_task * CHUNK_SIZE;  
            t.end = ((next_task+1)*CHUNK_SIZE < N ? (next_task+1)*CHUNK_SIZE : N);  
            MPI_Send(&t, sizeof(t), MPI_BYTE,  
                    worker, TAG_TASK, MPI_COMM_WORLD);  
            next_task++;  
        }  
  
        // 2) Loop: receive results, send new tasks  
        for (long completed = 0; completed < num_tasks; completed++) {  
            result_t r;  
            MPI_Status st;
```



```

    MPI_Recv(&r, sizeof(r), MPI_BYTE,
             MPI_ANY_SOURCE, TAG_RESULT,
             MPI_COMM_WORLD, &st);
    total += r.partial_sum;

    int worker = st.MPI_SOURCE;
    // If there's more work, send the next task:
    if (next_task < num_tasks) {
        task_t t;
        t.start = next_task * CHUNK_SIZE;
        t.end = ((next_task+1)*CHUNK_SIZE < N ? (next_task+1)*CHUNK_SIZE : N);
        MPI_Send(&t, sizeof(t), MPI_BYTE,
                 worker, TAG_TASK, MPI_COMM_WORLD);
        next_task++;
    }
    else {
        // Tell that worker to stop
        MPI_Send(NULL, 0, MPI_BYTE,
                 worker, TAG_STOP, MPI_COMM_WORLD);
    }
}

printf("Final sum = %.6f\n", total);
}
else {
    // ----- Worker -----
    while (1) {
        MPI_Status st;
        // Probe for an incoming task or stop tag
        MPI_Probe(0, MPI_ANY_TAG, MPI_COMM_WORLD, &st);
        if (st.MPI_TAG == TAG_STOP) {
            // Clean up and exit
            MPI_Recv(NULL, 0, MPI_BYTE, 0, TAG_STOP, MPI_COMM_WORLD,
MPI_STATUS_IGNORE);
            break;
        }
        // Otherwise it's a real task
        task_t t;
        MPI_Recv(&t, sizeof(t), MPI_BYTE, 0, TAG_TASK, MPI_COMM_WORLD,
MPI_STATUS_IGNORE);

        // Do the work
        result_t r;
        r.partial_sum = 0.0;
        r.start = t.start;
        r.end = t.end;
        for (long i = t.start; i < t.end; i++) {
            r.partial_sum += f(i);
        }

        // Send result back
        MPI_Send(&r, sizeof(r), MPI_BYTE,
                 0, TAG_RESULT, MPI_COMM_WORLD);
    }
}

MPI_Finalize();
return 0;
}

```

```
mpicc -O2 -o mpi_master_worker mpi_master_worker.c
mpirun -np 8 ./mpi_master_worker
```

A4

1)
Server

```
/* server.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define PORT 12345
#define BACKLOG 5
#define BUF_SIZE 64

/* Compute n! (note: long long overflows above ~20) */
static long long factorial(int n) {
    long long f = 1;
    int i;
    for (i = 2; i <= n; i++) {
        f *= i;
    }
    return f;
}

int main(void) {
    int listen_fd, conn_fd, opt, n;
    struct sockaddr_in srv_addr, cli_addr;
    socklen_t cli_len;
    uint32_t net_n;
    ssize_t r;
    char buf[BUF_SIZE];
    int blen;
    long long result;

    /* 1) create listening socket */
    listen_fd = socket(AF_INET, SOCK_STREAM, 0);
    if (listen_fd < 0) {
        perror("socket");
        exit(1);
    }

    /* allow address reuse */
    opt = 1;
    setsockopt(listen_fd, SOL_SOCKET, SO_REUSEADDR, &opt, sizeof(opt));

    /* bind to all interfaces */
    memset(&srv_addr, 0, sizeof(srv_addr));
    srv_addr.sin_family = AF_INET;
    srv_addr.sin_addr.s_addr = htonl(INADDR_ANY);
    srv_addr.sin_port = htons(PORT);

    if (bind(listen_fd, (struct sockaddr*)&srv_addr, sizeof(srv_addr)) < 0) {
        perror("bind");
    }
}
```

```

    close(listen_fd);
    exit(1);
}

if (listen(listen_fd, BACKLOG) < 0) {
    perror("listen");
    close(listen_fd);
    exit(1);
}

printf("Server listening on port %d\n", PORT);

for (;;) {
    cli_len = sizeof(cli_addr);
    conn_fd = accept(listen_fd, (struct sockaddr*)&cli_addr, &cli_len);
    if (conn_fd < 0) {
        perror("accept");
        continue;
    }

    /* 2) receive 32-bit integer n */
    r = recv(conn_fd, &net_n, sizeof(net_n), MSG_WAITALL);
    if (r != sizeof(net_n)) {
        fprintf(stderr, "recv failed\n");
        close(conn_fd);
        continue;
    }
    n = (int)ntohl(net_n);
    printf("Received n=%d from %s:%d\n",
        n,
        inet_ntoa(cli_addr.sin_addr),
        ntohs(cli_addr.sin_port));

    /* 3) compute factorial */
    result = factorial(n);

    /* 4) format result as ASCII string */
    blen = snprintf(buf, BUF_SIZE, "%lld", result);
    if (blen < 0 || blen >= BUF_SIZE) {
        fprintf(stderr, "snprintf error\n");
        close(conn_fd);
        continue;
    }

    /* 5) send NUL-terminated string */
    send(conn_fd, buf, blen+1, 0);

    close(conn_fd);
}

/* unreachable */
close(listen_fd);
return 0;
}

```

client

```

/* client.c */
#include <stdio.h>
#include <stdlib.h>

```

```

#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define PORT    12345
#define BUF_SIZE 64

int main(int argc, char *argv[]) {
    int sock, n;
    struct sockaddr_in srv_addr;
    char buf[BUF_SIZE];
    ssize_t r;
    uint32_t net_n;

    if (argc != 3) {
        fprintf(stderr, "Usage: %s <server-ip> <n>\n", argv[0]);
        return 1;
    }
    n = atoi(argv[2]);

    /* 1) create socket */
    sock = socket(AF_INET, SOCK_STREAM, 0);
    if (sock < 0) {
        perror("socket");
        return 1;
    }

    /* 2) configure server address */
    memset(&srv_addr, 0, sizeof(srv_addr));
    srv_addr.sin_family = AF_INET;
    srv_addr.sin_port = htons(PORT);
    if (inet_aton(argv[1], &srv_addr.sin_addr) == 0) {
        fprintf(stderr, "Invalid IP: %s\n", argv[1]);
        close(sock);
        return 1;
    }

    /* 3) connect */
    if (connect(sock, (struct sockaddr*)&srv_addr, sizeof(srv_addr)) < 0) {
        perror("connect");
        close(sock);
        return 1;
    }

    /* 4) send n in network byte order */
    net_n = htonl((uint32_t)n);
    if (send(sock, &net_n, sizeof(net_n), 0) != sizeof(net_n)) {
        perror("send");
        close(sock);
        return 1;
    }

    /* 5) receive NUL-terminated ASCII factorial */
    r = recv(sock, buf, BUF_SIZE, MSG_WAITALL);
    if (r <= 0) {
        fprintf(stderr, "recv failed\n");
        close(sock);
        return 1;
    }
}

```

```

    buf[BUF_SIZE-1] = '\0'; /* ensure NUL termination */

    printf("Factorial of %d is %s\n", n, buf);
    close(sock);
    return 0;
}

```

2)
mapper.c

```

/* mapper.c */
#include <stdio.h>
#include <stdlib.h>
#include <dirent.h>
#include <sys/stat.h>
#include <pwd.h>
#include <string.h>
#include <unistd.h>

int main(void) {
    DIR *d = opendir(".");
    if (!d) {
        perror("opendir");
        return 1;
    }

    struct dirent *ent;
    struct stat st;
    while ((ent = readdir(d)) != NULL) {
        /* skip "." and ".." */
        if (strcmp(ent->d_name, ".") == 0 || strcmp(ent->d_name, "..") == 0)
            continue;

        if (stat(ent->d_name, &st) < 0) {
            /* could not stat—skip */
            continue;
        }
        /* only regular files */
        if (!S_ISREG(st.st_mode))
            continue;

        /* lookup owner name */
        struct passwd *pw = getpwuid(st.st_uid);
        const char *user = pw ? pw->pw_name : "UNKNOWN";

        /* emit: size \t user */
        printf("%lld\t%s\n", (long long)st.st_size, user);
    }

    closedir(d);
    return 0;
}

gcc -O2 -o mapper mapper.c

```

reducer.c

```

/* reducer.c */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

```

```

#define MAX_USERS 1024
#define MAX_USER_LEN 256

int main(void) {
    char line[1024];
    long long max_size = -1;
    char users[MAX_USERS][MAX_USER_LEN];
    int user_count = 0;

    while (fgets(line, sizeof(line), stdin)) {
        long long sz;
        char user[MAX_USER_LEN];
        if (sscanf(line, "%lld\t%255s", &sz, user) != 2)
            continue;

        if (sz > max_size) {
            /* new maximum: reset list */
            max_size = sz;
            user_count = 0;
            strncpy(users[user_count++], user, MAX_USER_LEN-1);
            users[user_count-1][MAX_USER_LEN-1] = '\0';
        }
        else if (sz == max_size) {
            /* add if not already in list */
            int found = 0;
            for (int i = 0; i < user_count; i++) {
                if (strcmp(users[i], user) == 0) {
                    found = 1;
                    break;
                }
            }
            if (!found && user_count < MAX_USERS) {
                strncpy(users[user_count++], user, MAX_USER_LEN-1);
                users[user_count-1][MAX_USER_LEN-1] = '\0';
            }
        }
    }

    /* output the owners of largest files */
    for (int i = 0; i < user_count; i++) {
        printf("%s\n", users[i]);
    }
    return 0;
}

gcc -O2 -o reducer reducer.c

```

A5

```

file_transfer.x
program FILE_TRANSFER_PROG {
    version FILE_TRANSFER_VERS {
        /* get_file: client->server filename (string), server->client file data (opaque) */
        getfile STRING -> FILEDATA = 1;
        /* put_file: client->server filename + data, server->client status (int) */
        putfile PUTARGS -> int = 2;
    } = 1;
} = 0x20000001;

/* Data types */
typedef string STRING<255>;

```

```
struct FILEDATA {
    unsigned long size;
    opaque data<>;
};
```

```
struct PUTARGS {
    STRING filename;
    unsigned long size;
    opaque data<>;
};
```

```
rpcgen -a file_transfer.x
```

```
file_transfer_svc.c
#include "file_transfer.h"
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <rpc/rpc.h>
```

```
/* Handler for getfile_1 */
FILEDATA *
getfile_1_svc(STRING *filename, struct svc_req *rqstp)
{
    static FILEDATA result;
    FILE *f;
    unsigned long n;

    f = fopen(filename->str, "rb");
    if (!f) {
        result.size = 0;
        result.data.data_len = 0;
        return &result;
    }
    fseek(f, 0, SEEK_END);
    n = ftell(f);
    rewind(f);

    result.size = n;
    result.data.data_len = n;
    result.data.data_val = malloc(n);
    fread(result.data.data_val, 1, n, f);
    fclose(f);
    return &result;
}
```

```
/* Handler for putfile_1 */
int *
putfile_1_svc(PUTARGS *args, struct svc_req *rqstp)
{
    static int status;
    FILE *f;

    f = fopen(args->filename.filename, "wb");
```

```

    if (!f) {
        status = -1;
        return &status;
    }
    fwrite(args->data.data_val, 1, args->size, f);
    fclose(f);
    status = 0;
    return &status;
}

```

```

file_transfer_clnt.c
#include "file_transfer.h"
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

```

```

void usage(const char *prog) {
    fprintf(stderr,
        "Usage: %s <server> get <filename> | put <localfile> <remotefile>\n",
        prog);
    exit(1);
}

```

```

int main(int argc, char **argv) {
    CLIENT *cl;
    char *host;

    if (argc < 4) usage(argv[0]);
    host = argv[1];
    cl = clnt_create(host, FILE_TRANSFER_PROG, FILE_TRANSFER_VERS, "tcp");
    if (!cl) {
        clnt_pcreateerror(host);
        exit(1);
    }

```

```

    if (strcmp(argv[2], "get") == 0) {
        STRING filename;
        FILEDATA *resp;
        filename.str = argv[3];
        filename.str_len = strlen(argv[3]);
        resp = getfile_1(&filename, cl);
        if (resp && resp->size > 0) {
            fwrite(resp->data.data_val, 1, resp->size, stdout);
        } else {
            fprintf(stderr, "Error: could not fetch '%s'\n", argv[3]);
        }
    }

```

```

}
else if (strcmp(argv[2], "put") == 0 && argc==5) {
    PUTARGS args;
    int *rstatus;
    FILE *f = fopen(argv[3], "rb");
    if (!f) { perror("fopen"); exit(1);}

```



```

fseek(f,0,SEEK_END); args.size = ftell(f); rewind(f);
args.data.data_len = args.size;
args.data.data_val = malloc(args.size);
fread(args.data.data_val,1,args.size,f);
fclose(f);
args.filename.filename = argv[4];
args.filename.filename_len = strlen(argv[4]);

rstatus = putfile_1(&args, cl);
if (!rstatus || *rstatus!=0)
    fprintf(stderr, "Error: putfile failed\n");
}
else {
    usage(argv[0]);
}

clnt_destroy(cl);
return 0;
}

```

A7

```

1)
a)
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define TAG_REQ  1
#define TAG_ACK  2
#define TAG_REL  3

int main(int argc, char **argv) {
    MPI_Init(&argc, &argv);
    int P, rank;
    MPI_Comm_size(MPI_COMM_WORLD, &P);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Lamport clock
    int clock = 0;
    // Highest timestamp of our own request (if pending), else -1
    int req_ts = -1;
    // Count of outstanding ACKs
    int pending_acks = 0;
    // Deferred requests queue: we only need to know which ranks to ACK later
    int deferred[128], defcnt = 0;

    // For demonstration, each process enters CS exactly once, after a barrier
    MPI_Barrier(MPI_COMM_WORLD);
    // 1) Request CS
    clock++;

```

```

req_ts = clock;
pending_acks = P-1;
// broadcast request(ts, rank)
for (int dst = 0; dst < P; dst++) if (dst != rank) {
    MPI_Send(&req_ts, 1, MPI_INT, dst, TAG_REQ, MPI_COMM_WORLD);
}
printf("[R%d] Requested CS at ts=%d, waiting %d acks\n", rank, req_ts, pending_acks);

// 2) Process incoming messages until we collect all ACKs
while (pending_acks > 0) {
    MPI_Status st;
    int in_ts;
    MPI_Recv(&in_ts, 1, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &st);
    clock = clock > in_ts ? clock+1 : in_ts+1;

    if (st.MPI_TAG == TAG_REQ) {
        int src = st.MPI_SOURCE;
        // if our own request is earlier, defer; else ack immediately
        if (req_ts >= 0 && (req_ts < in_ts || (req_ts == in_ts && rank < src))) {
            deferred[defcnt++] = src;
        } else {
            MPI_Send(&clock, 1, MPI_INT, src, TAG_ACK, MPI_COMM_WORLD);
        }
    }
    else if (st.MPI_TAG == TAG_ACK) {
        pending_acks--;
    }
}

// 3) ENTER Critical Section
printf("[R%d] ENTER CS at ts=%d\n", rank, clock);
sleep(1);           // simulate work
printf("[R%d] LEAVE CS at ts=%d\n", rank, clock);

// 4) RELEASE: send release to deferred
for (int i = 0; i < defcnt; i++) {
    int dst = deferred[i];
    MPI_Send(&clock, 1, MPI_INT, dst, TAG_ACK, MPI_COMM_WORLD);
}

MPI_Finalize();
return 0;
}
mpicc -O2 -o lamport lamport.c
mpirun -np 5 ./lamport

```

b)

```

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

```

```

#define TAG_REQ 1
#define TAG_REPLY 2

int main(int argc, char **argv) {
    MPI_Init(&argc, &argv);
    int P, rank;
    MPI_Comm_size(MPI_COMM_WORLD, &P);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    int clock = 0;
    int req_ts = -1, replies_needed = 0;
    int deferred[128], defcnt = 0;

    MPI_Barrier(MPI_COMM_WORLD);

    // 1) send REQUEST to all
    clock++;
    req_ts = clock;
    replies_needed = P-1;
    for (int dst = 0; dst < P; dst++) if (dst != rank) {
        MPI_Send(&req_ts, 1, MPI_INT, dst, TAG_REQ, MPI_COMM_WORLD);
    }
    printf("[R%d] Req CS ts=%d\n", rank, req_ts);

    // 2) wait for replies
    while (replies_needed > 0) {
        MPI_Status st;
        int in_ts;
        MPI_Recv(&in_ts, 1, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &st);
        clock = clock > in_ts ? clock+1 : in_ts+1;

        if (st.MPI_TAG == TAG_REQ) {
            int src = st.MPI_SOURCE;
            // defer if our request is earlier, else reply immediately
            if (req_ts >= 0 && (req_ts < in_ts || (req_ts == in_ts && rank < src))) {
                deferred[defcnt++] = src;
            } else {
                MPI_Send(&clock, 1, MPI_INT, src, TAG_REPLY, MPI_COMM_WORLD);
            }
        }
        else if (st.MPI_TAG == TAG_REPLY) {
            replies_needed--;
        }
    }

    // 3) Critical Section
    printf("[R%d] ENTER CS ts=%d\n", rank, clock);
    sleep(1);
    printf("[R%d] LEAVE CS ts=%d\n", rank, clock);
}

```

```

// 4) reply to deferred
for (int i = 0; i < defcnt; i++) {
    MPI_Send(&clock, 1, MPI_INT, deferred[i], TAG_REPLY, MPI_COMM_WORLD);
}

```

```

MPI_Finalize();
return 0;
}

```

```

mpicc -O2 -o ricart ricart.c
mpirun -np 5 ./ricart

```

2)

```

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

```

```

#define TAG_REPORT 1
#define TAG_DETECT 2

```

```

// Max processes and edges
#define MAXP 64
#define MAXE 128

```

```

// A single edge: process 'u' is waiting for process 'v'
typedef struct {
    int u, v;
} edge_t;

```

```

// Each worker reports its edges in a variable-length array
typedef struct {
    int count;
    edge_t edges[MAXE];
} wfg_t;

```

```

// coordinator: detect cycle in directed graph
int has_cycle(int P, int adj[][MAXP]) {
    int visited[MAXP] = {0}, stack[MAXP] = {0};
    // DFS from each node
    for (int i = 1; i < P; i++) {
        memset(visited, 0, sizeof(visited));
        memset(stack, 0, sizeof(stack));
        // inner recursive lambda simulated by explicit stack:
        int st_sz = 0;
        // push (i, next_child=1)
        int node = i, next = 1;
        int call_stack[MAXP][2];
        call_stack[st_sz][0] = node; call_stack[st_sz][1] = next;
        st_sz++; visited[node] = 1; stack[node] = 1;
        while (st_sz) {
            node = call_stack[st_sz-1][0];

```

```

    next = call_stack[st_sz-1][1];
    if (next >= P) {
        // done with node
        stack[node] = 0;
        st_sz--;
        continue;
    }
    call_stack[st_sz-1][1]++; // next time, try child+1
    if (adj[node][next]) {
        if (stack[next]) return 1;
        if (!visited[next]) {
            // push child
            visited[next] = 1; stack[next] = 1;
            call_stack[st_sz][0] = next;
            call_stack[st_sz][1] = 1;
            st_sz++;
        }
    }
}
}
return 0;
}

```

```

int main(int argc, char **argv) {
    MPI_Init(&argc, &argv);
    int P, rank;
    MPI_Comm_size(MPI_COMM_WORLD, &P);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if (rank == 0) {
        // coordinator
        int adj[MAXP][MAXP];
        memset(adj, 0, sizeof(adj));
        wfg_t report;
        MPI_Status st;

        while (1) {
            // 1) ask all workers to report
            for (int src = 1; src < P; src++) {
                MPI_Send(NULL, 0, MPI_BYTE, src, TAG_DETECT, MPI_COMM_WORLD);
            }
            // 2) gather
            memset(adj, 0, sizeof(adj));
            for (int src = 1; src < P; src++) {
                MPI_Recv(&report, sizeof(report), MPI_BYTE, src,
                    TAG_REPORT, MPI_COMM_WORLD, &st);
                for (int e = 0; e < report.count; e++) {
                    adj[ report.edges[e].u ][ report.edges[e].v ] = 1;
                }
            }
            // 3) detect
            if (has_cycle(P, adj)) {

```

```

        printf("[COORD] Deadlock detected!\n");
    } else {
        printf("[COORD] No deadlock.\n");
    }
    fflush(stdout);
    sleep(5); // periodic
}
}
else {
    // worker: simulate holding/waiting
    wfg_t local;
    MPI_Status st;

    // for demo, each worker cycles through some fixed edges
    while (1) {
        // wait for detect request
        MPI_Recv(NULL, 0, MPI_BYTE, 0, TAG_DETECT, MPI_COMM_WORLD, &st);

        // build local wait-for graph edges
        // e.g. rank i waits for (i%P)+1
        local.count = 1;
        local.edges[0].u = rank;
        local.edges[0].v = (rank % (P-1)) + 1;

        // send report
        MPI_Send(&local, sizeof(local), MPI_BYTE, 0,
                 TAG_REPORT, MPI_COMM_WORLD);
    }
}

MPI_Finalize();
return 0;
}
mpicc -O2 -o deadlock deadlock.c
# Centralized deadlock detection with 5 ranks (1 coordinator + 4 workers):
mpirun -np 5 ./deadlock

```

A